

training loss (SGD)

- too high (3.0): overshoots, settles worse
- sensible (0.3)
- too low (0.00001): asleep

